

CMPT 506

Database Concurrency Control



Read Chapters 20 and 21

Dr. Abdelkarim Erradi

Department of Computer Science and Engineering

QU

Outline

- DB Transaction as a mechanism to achieve Atomicity
- Concurrency control
- Lock-based Concurrency control
- Transaction Isolation Levels
- Time-based Concurrency control
- Optimistic Concurrency Control


DB Transaction as a mechanism to achieve Atomicity

Transactions **ACID** properties

- **Atomic:** Everything in a transaction succeeds or the entire transaction is rolled back (All or Nothing)
- **Consistent:** data affected meet all validation rules such as constraints
- **Isolated:** Transactions cannot interfere with each other => The updates of a transaction must not be made visible to other transactions until it is **committed**
- **Durable:** Results from completed transactions survive failures

Atomicity

- Consider a bank transaction T:

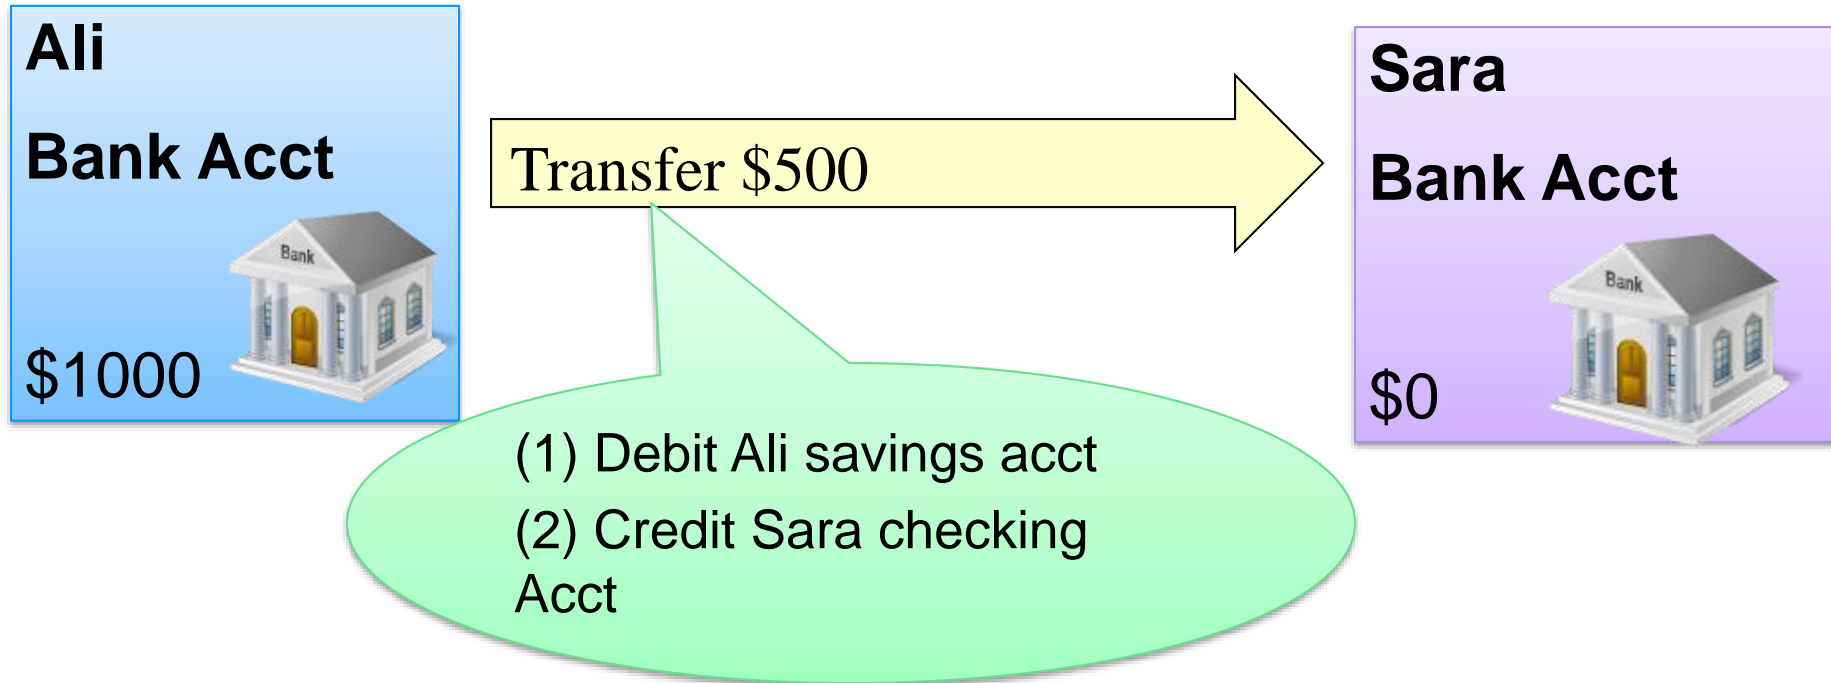


- T is transferring \$100 from B's account to A's account.
 - What if there is an error right after the first statement of T has been executed, i.e. the second statement is not executed?
- => You will get **partial update** = inconsistent state
- The DBS has to ensure that every transaction is treated as an *atomic* unit, i.e. either all succeed otherwise rollback

Concept of Transaction

- **Atomicity is achieved using Transaction**
- Transaction = Logical unit of work on the database
 - Transfer money from one bank account to another -
 - Checkout: place order and process payment
- A transaction consists of a sequence of read / write operations that must be performed as a **single** logical unit
- Transaction boundaries are defined by the database user / application programmer

A transaction is a sequence of operations that must be executed as a whole



Either both (1) and (2) happen or neither!

Every DB action takes place inside a transaction

We abstract away most of the application code when thinking about transactions

User's point of view

Transfer \$500

1. Debit savings
2. Credit checking

Programmer's point of view

Read Balance1
Write Balance1
Read Balance2
Write Balance2

DB's point of view

Transaction = a sequence of DB reads (R) and writes (W)

T: R(A), W(A), R(B), W(B)

→ time

SQL Transaction

- By default, each SQL statement (any query or modification of the database or its schema) is treated as a separate transaction
- Transactions can also be defined explicitly

```
START TRANSACTION;  
    <sequence of SQL statements>  
COMMIT;  
or ROLLBACK;
```
- COMMIT makes all modifications of the transaction permanent, ROLLBACK undoes all DB modifications made

Concurrency control

Concurrency Control

- Multiple **concurrent** transactions T_1, T_2, \dots may read/write Data Items A_1, A_2, \dots concurrently
- Concurrency Control is the process of managing concurrent operations performed on shared data so that data manipulation does not generate inconsistent databases or produce wrong results
- Objective
 - Maximise throughput (i.e., work performed)
 - Minimize response time
- Constraint
 - Avoid **interference** between transactions

Three Concurrency Anomalies

- **Lost update**

- Two transactions T_1 and T_2 both modify the same data
- T_1 and T_2 both commit
- Final state shows effects of only T_1 but not of T_2

- **Dirty read**

- T_1 reads data written by T_2 while T_2 has not committed
- If T_2 aborts then T_1 will have dirty data

- **Unrepeatable read**

- Getting inconsistent results when a read operation is reexecuted within a Transaction T

Illustrative Example

- **Example (to illustrate consistency issues that can be introduced by concurrent updates)**
- Ali at ATM1 withdraws \$100
- Sara at ATM2 withdraws \$50
- Initial balance = \$400, final balance = ?
 - Should be \$250 no matter who goes first

Read balance from DB;

If balance > withdrawalAmount {
 balance = balance - withdrawalAmount;
 Write balance to DB;

}

No concurrent transactions scenario

Ali withdraws \$100:

```
read balance; ==> $400
if balance > amount then
    balance = balance - amount; ==> $300
write balance; ==> $300
```

Sara withdraws \$50:

```
read balance; ==> $300
if balance > amount then
    balance = balance - amount; ==> $250
write balance; ==> $250
```

Lost update problem

Ali withdraws \$100:

```
read balance; => $400
```

Sara withdraws \$50:

```
read balance; => $400
```

```
If balance > amount then
```

```
    balance = balance - amount; => $350
```

```
write balance; => $350
```

```
if balance > amount then
```

```
    balance = balance - amount; => $300
```

```
write balance; => $300
```



Lost update problem => DB is in inconsistent state

Lost update problem

Ali withdraws \$100:

```
read balance; => $400
```

```
if balance > amount then  
  balance = balance - amount; => $300  
write balance; => $300
```

Sara withdraws \$50:

```
read balance; => $400
```

```
if balance > amount then  
  balance = balance - amount; => $350  
write balance; => $350
```




Lost update problem => DB is in inconsistent state

Dirty read problem

What will be the final account balance?



Transaction 1:

Add \$100 to
account A



Transaction 2:

Add \$200 to
account A

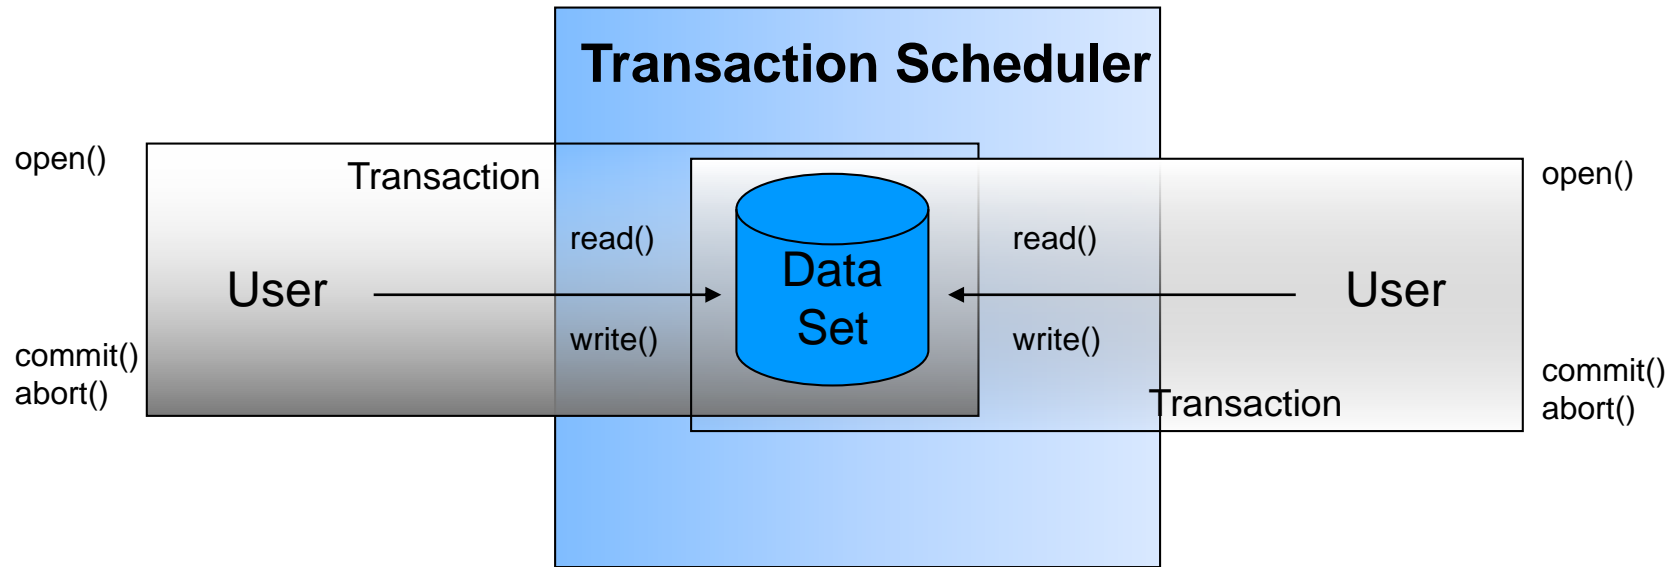


Dirty read problem

Isolation

- To preserve consistency the DBMS must manage concurrency to **guarantee the Isolation property**.
 - Independence from all other transactions (serializability) => **same results as if the statements would have been executed in a single user scenario**
 - A transaction should not be affected by other concurrently running transactions


Scheduling Concurrent Transactions



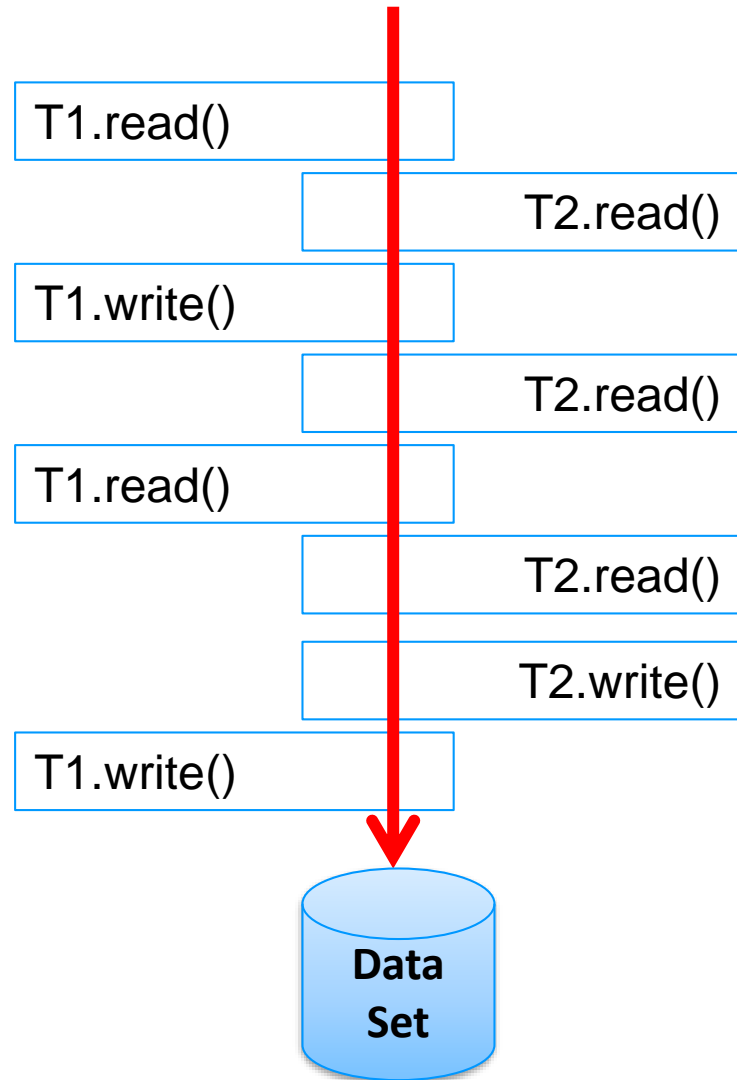
- The transaction scheduler has to organise or “schedule” database read and write actions of all concurrent transactions
- A “**schedule**” is a sequence of operations from different transactions that may be executed in an interleaved fashion
 - Such a schedule may compromise the integrity / consistency of a database

Serial vs. Non-serial Schedule

Serial Schedule



Non-serial / Interleaved Schedule



Serial schedule in which T_1 precedes T_2

T_1	T_2	A	B
		25	25
READ(A,t)			
t := t+100			
WRITE(A,t)		125	
READ(B,t)			
t := t+100			
WRITE(B,t)			125
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		250

Constraint of $A = B$ is satisfied after this schedule

Serial schedule in which T_2 precedes T_1

T_1	T_2	A	B
		25	25
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	50	
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		50
READ(A,t)			
t := t+100			
WRITE(A,t)		150	
READ(B,t)			
t := t+100			
WRITE(B,t)			150

Constraint of $A = B$ is satisfied after this schedule

Interleaved Transaction Schedules

- We want that the database system schedules transactions in an interleaved fashion:
 - Improve the responsiveness and increase throughput
- Interleaved schedules also create problems
 - Transactions may overwrite each others' updates
 - Transactions may base their calculations on retrieved data that is already out-of-date or on “dirty reads”

Serializable Transaction Schedules

- In order to avoid the concurrency problems described, one obvious solution would be to **schedule only one transaction at a time** for execution
- Such a completely “serialised” schedule will ensure that the **transactions are completely isolated** and cannot interfere with each other
 - But this strategy will reduce **concurrency** and **throughput**

Serializable Schedule

- DB will try to find a non-serial schedule that is ***equivalent*** to a serial schedule:
 - Schedule is **serializable** if its effect on the database state is the same as that of some serial schedule
- A **serializable schedule** is used by the Transaction Manager to schedule operations of different transactions in a way that interference and problems such as “lost updates” are avoided

A serializable, but not serial, schedule

T_1	T_2	A	B
		25	25
READ(A,t)			
t := t+100			
WRITE(A,t)		125	
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
READ(B,t)			
t := t+100			
WRITE(B,t)			125
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		250

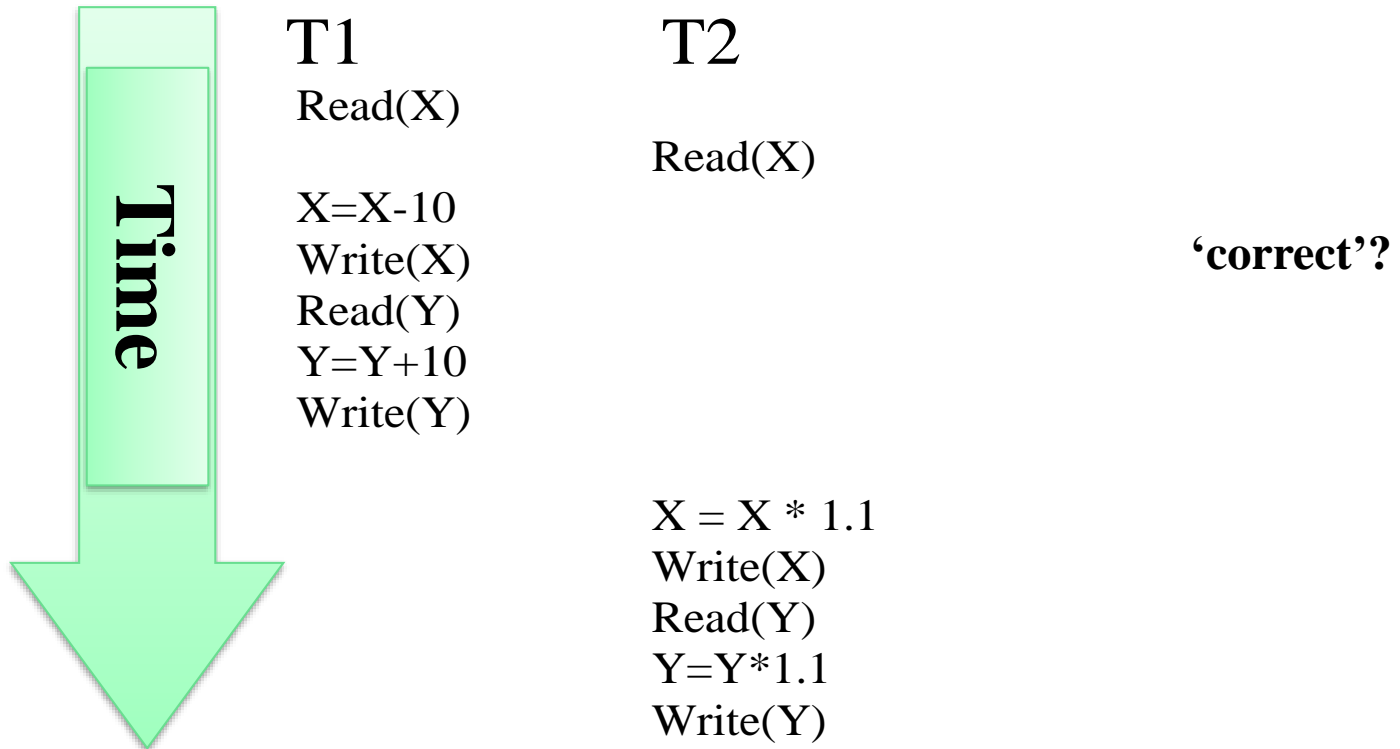
Constraint of $A = B$ is satisfied after this schedule

Non-serializable schedule

T_1	T_2	A	B
		25	25
READ(A,t)			
t := t+100			
WRITE(A,t)		125	
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		50
READ(B,t)			
t := t+100			
WRITE(B,t)			150

Constraint of $A = B$ is violated after this schedule => Non-serializable schedule

Interleaved execution



Schedule: The order of execution of operations of two or more transactions

How to define correctness?

Let's start from something definitely correct:

Serial executions

T1

Read(X)

X=X-10

Write(X)

Read(Y)

Y=Y+10

Write(Y)

T2

‘correct’

by definition

Read(X)

X = X * 1.1

Write(X)

Read(Y)

Y=Y*1.1

Write(Y)


A schedule is ‘correct’ if it is **serializable**
i.e., equivalent to a serial schedule

How to draw a Precedence Graph?

- Precedence Graph = {**Nodes**: transactions,
Arcs: r/w or w/w conflicts}
- The precedence graph for a schedule S contains:
 - A node for each committed transaction in S
 - An **arc** from T_i to T_j if an action of T_i **precedes and conflicts with** one of T_j 's actions.
- The **schedule S is serializable if and only if the precedence graph has no cycles.**

Example 1

T1	T2	T3
Read(A)		
...		
Write(A)		
		Read(A)
		...
		Write(A)
	Read(B)	
	...	
	Write(B)	
Read(B)		
...		
Write(B)		



serial execution?


Serializable Schedule: **T2, T1, T3**

(Notice that T3 should go after T2, although it starts before it!) => always 'read before write'

Precedence Graph Example 2

$$D = \begin{bmatrix} T1 & T2 & T3 \\ R(A) & & \\ & W(A) & \\ W(A) & & \\ & & W(A) \end{bmatrix}$$

$$D = R_1(A) \ W_2(A) \ W_1(A) \ W_3(A)$$



As T_1 and T_2 constitute a cycle the above schedule is not serializable.

Example 3 - 'Lost-update' problem

T1
Read(N)

$N = N - 1$


Write(N)

T2

Read(N)

$N = N - 1$

Write(N)



Cycle -> not serializable

Not equivalent to any serial execution (why not?) -> incorrect!

We can draw a precedence graph to prove this

Concurrency Control Protocols

- The basis of concurrency control is protocols to maintain serialization in DBMSs
- E.g. protocols
 - **Two-Phase Locking (2PL)**
 - **Timestamps Ordering**
 - **Optimistic Concurrency Control (OCC)**
 - **Etc.**

Lock-based Concurrency control

Enforcing Serializability by Locks


How to achieve correct concurrency control?

- **Locks** = most popular solution for concurrency control
- Lock manager: grants/denies lock requests
- Locking mechanisms prevent conflicts
 - Readers block writers
 - Writers block readers

How lock works

- Transaction needs lock before read or write
- Transaction must ask for lock
- If it does not get the lock
 - maybe another transaction already holds the lock
 - it is suspended
 - put on queue waiting for lock
- When transaction releases lock
 - Some suspended transaction awakened and given lock.

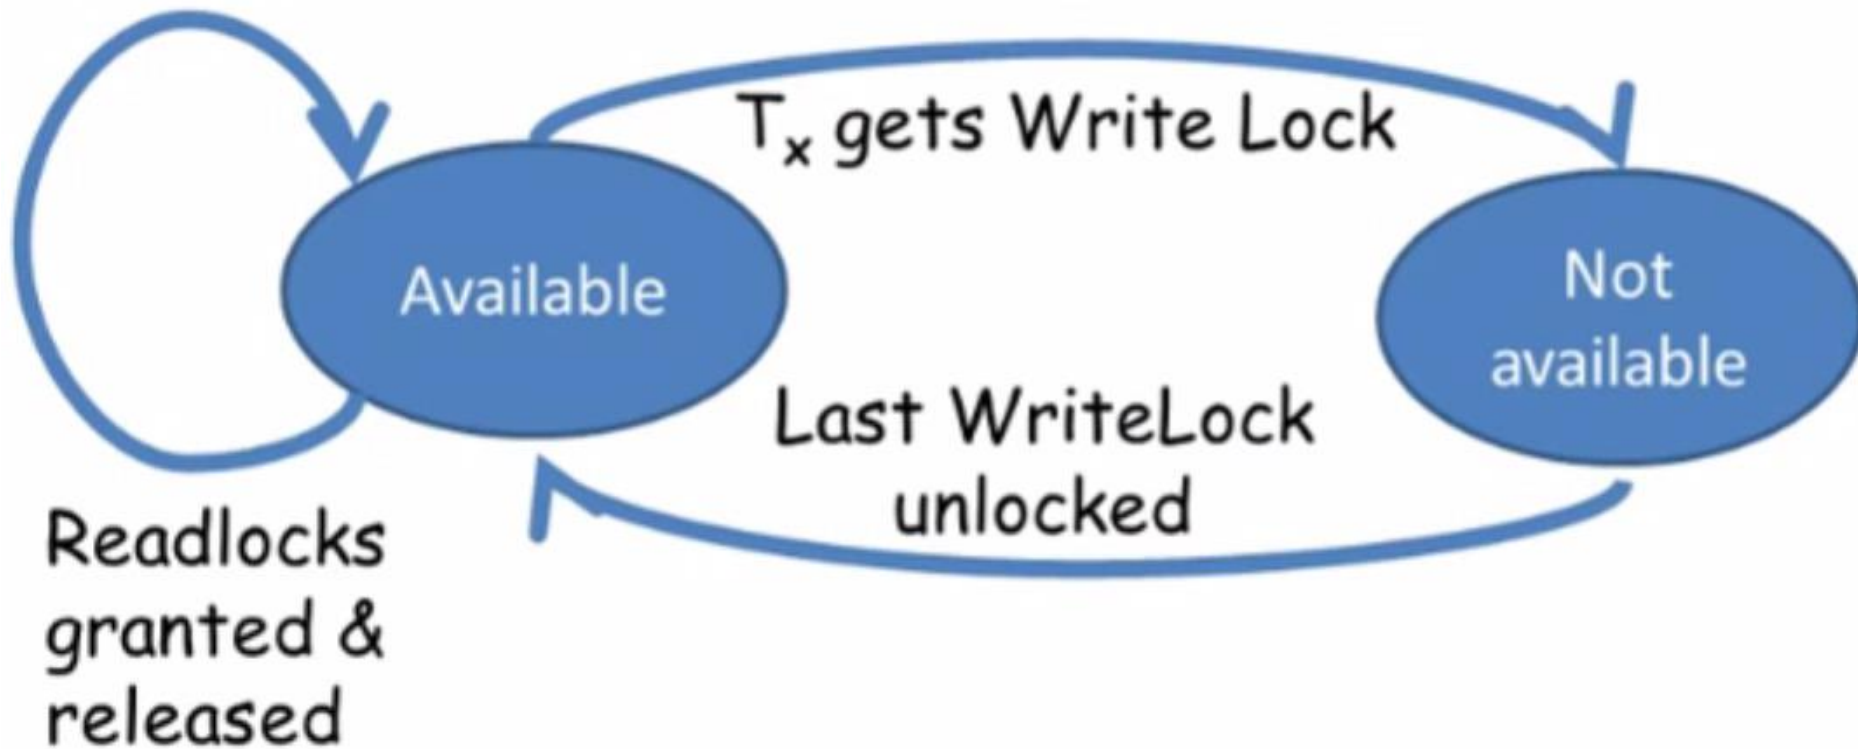
Lock State Diagram



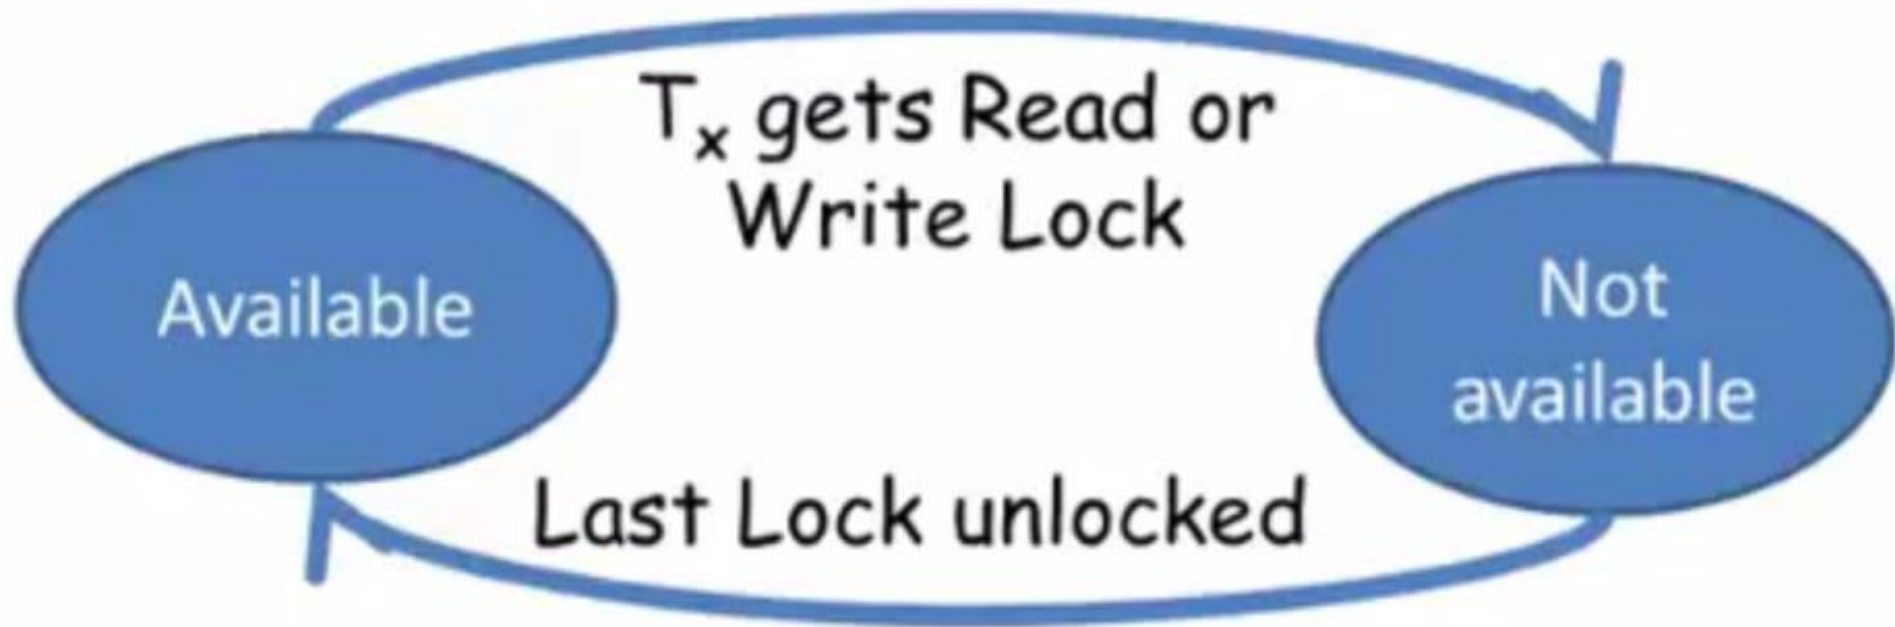
Multi-transaction Rules

- Multiple Read Locks OK
 - Called Shared Lock
 - Several transactions can read same info
- Write Lock must be only lock on item
 - 2 Write Locks lead to race condition: lost update
 - Sharing with Read lock leads to dirty read or non-repeatable read.

Read Lock State Diagram



Write Lock State Diagram



ReadLock WriteLock Compatibility Matrix

First
Transaction
Holds

Read Lock

Write Lock

Second Transaction Wants

Read Lock



Write Lock



Locking protocol

- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks
- Most popular protocol:
 - **Strict 2 Phase Locking (2PL)** = a transaction must hold all its exclusive locks till it commits/aborts
- **THEOREM**: if all transactions obey 2PL -> all schedules are serializable
 - But reduces concurrency


2PL

- Phase 1: **Growing Phase**
 - transaction may obtain locks
 - transaction may not release locks
- Phase 2: **Shrinking Phase**
 - transactions issue no lock/upgrade request, after the first unlock/downgrade
 - transaction may release locks
 - transaction may not obtain locks
- 2PL assures serializability

Deadlocks

- The use of locks may lead to deadlocks
- **Definition**
 - A deadlock is a state in which each member of a group of transactions is waiting for some other member to release a lock
- The scheduler is responsible for detecting and breaking the deadlock
 - Deadlocks may be broken by simply aborting one of the transactions involved
- How do you choose which transaction to abort?
 - Abort the oldest
 - Abort depending on the complexity of the transactions

Deadlocks



- Auto-detection and “resolution”

Msg 1205, Level 13, State 51, Line 1

Transaction (Process ID 232) was deadlocked on lock resources with another process and has been chosen as the deadlock victim. Rerun the transaction

Deadlock Example

Transaction A

Read X **Writelock X**

Read Y **Readlock Y**

Write X

Cannot get lock
due to B's
Writelock -- A
sleeps

Each sleeps, waiting for the other

Transaction B

Writelock Y
Read Y

Readlock X
Read X

Write Y

Cannot get
lock due to
A's Writelock
-- B sleeps

Handling Deadlocks

- Deadlock Detection
 - Use the Precedence graph to identify cycles and select transactions to be aborted
 - Select the transaction involved in most of the cycles
 - Select the oldest transactions
 - Select the one that did the least amount of work
- Deadlock Prevention
 - Lock all objects at the very beginning of a transaction in one atomic action
 - Problem
 - Reduced concurrency: unnecessary access restriction to shared resources

Handling Deadlocks using Timeout

- Rather than detecting deadlock (an overhead), you could just use timeouts, but how long should the timeout be?
- Timeouts
 - Each lock is given a period of time where it is invulnerable
 - After this timeout, it becomes vulnerable
 - If transaction X holds a lock that becomes vulnerable and transaction Y is waiting for X, then X is aborted
 - Problem
 - Hard to decide on an appropriate length of timeout
 - Transactions may be aborted, when the lock becomes vulnerable and another transaction waits, but there is no deadlock

Transaction Isolation Levels

Transaction Isolation Levels

- Determine how much one user impacts another
- 4 standard levels
 - Read Uncommitted
 - Read Committed (default in many DMBS)
 - Repeatable Read
 - Serializable
- The *isolation level* of a transaction defines what data *that* transaction may see

Read Committed (the default)

- Cannot read uncommitted data
- Share lock before reading data and released after processing is complete
- Dirty reads do not happen but non-repeatable reads can happen

Read Uncommitted

- Does not take or honour locks
- Can read uncommitted data
 - *Dirty data* is data that has been modified by a transaction that has not yet committed.
 - If that transaction is rolled back after another transaction has read its dirty data, inconsistency is introduced
- Sacrificing consistency in favour of high concurrency
- Useful for reporting applications
- Be very careful!

Repeatable Read

- All **data** records read by a SELECT statement cannot be changed
- Holds shared locks on data read until transaction commit/rollback
- Rows that have been read can be read again with confidence they won't have changed
 - A query running more than once within the same transaction returns same values
- If the SELECT statement contains any ranged WHERE clauses, phantom reads can occur
- Reduce concurrency and degrade performance

Unrepeatable Reads Example

- A transaction reads the same data item twice, with different results!
 - Start a transaction (T1) under repeatable read isolation level. Query the *ACCOUNTS* table with predicate (*account_balance* > 1000). Let us say it returns 10 rows
 - Another transaction (T2) cannot update *ACCOUNTS* with *account_balance* > 1000
 - The reason is that transaction T1 only locked the 10 qualifying rows but did not lock the predicate range. With the result, the transaction T2 could insert a new row in the same predicate range.

Phantom Problem

T1

T2

```
SELECT *  
FROM Product  
WHERE color='blue'
```

```
INSERT INTO Product(name, color)  
VALUES ('Shrek Toy','blue')
```

```
SELECT *  
FROM Product  
WHERE color='blue'
```


Phantom Problem

- A “phantom” is a tuple that is invisible during part of a transaction execution but appears when the query is re-executed
- In our example:
 - T1: reads list of products
 - T2: inserts a new product
 - T1: re-reads: a new product appears !

Dealing With Phantoms

- Lock the entire table, or
- Lock the index entries of qualifying rows
 - if index is available

Dealing with phantoms is expensive !

Serializable

- This isolation level specifies that all transactions occur in a completely isolated fashion
 - *i.e.*, as if all transactions in the system had executed serially, one after the other
- Locks index ranges as well as rows or table locks
- Phantom rows will not appear if the same query is issued twice within a transaction
- Greatly reduces concurrency

Concurrency and Consistency

Isolation Table	Dirty Read	Non-Repeatable Read	Phantom Read
Read Uncommitted	Possible	Possible	Possible
Read Committed	Not Possible	Possible	Possible
Repeatable Read	Not Possible	Not Possible	Possible
Serializable	Not Possible	Not Possible	Not Possible

- Concurrency and consistency are mutually opposing goals

Time-based Concurrency control

Timestamp Ordering

- The timestamp of a transaction T is the time at which that transaction was initiated in the DBMS: $TS(T)$
- We can use clock time or an incremental identifier (counter) for $TS(T)$
- Two timestamps are also associated with each data item x .
 1. $read_TS(x)$ is the $TS(T)$ of the last transaction T to read from x .
 2. $write_TS(x)$ is the $TS(T)$ of the last transaction T to write to x .

Protocol Rules

- Two simple rules to follow:
 1. Before T issues a write(x), check to see if
 - $TS(T) < read_TS(x)$ or if
 - $TS(T) < write_TS(x)$
 - If so, then abort transaction T.
 - If not, then perform write(x) and set $write_TS(x) = TS(T)$
 2. Before T issues a read(x), check to see if
 - $TS(T) < write_TS(x)$ Then abort transaction T.
 - if $TS(T) \geq write_TS(x)$ then execute read(x) and set $read_TS(x) = TS(T)$ only if TS(T) is greater than the current $read_TS(x)$

Advantages and Limitation

- When a transaction is aborted, it is the restarted and issued a new $TS(T)$.
- Note that with timestamp ordering, deadlock can not occur.
- However, starvation is possible i.e., a transaction keeps getting aborted over and over.

Cascading rollbacks

- Timestamp ordering can also produce cascading rollbacks:
 - Assume transaction T begins executing and performs some read and write operations on data items *a*, *b* and *c*
 - However, T then reaches a data item it can not read or write and T must then be aborted.
 - Any effects of transaction T must then be rolled back.
 - Before T aborts, however, other transactions (T1, T2 and T3) have read and written data items *a*, *b* and *c* so *these other transactions must also be rolled back*.
 - There may be other transactions (T4 and T5) that worked with data items read or written by T1, T2 and T3, etc.

Optimistic Concurrency Control

Optimistic Concurrency Control

- Two Phase Locking (2PL) and Timestamp Ordering (TO) are *pessimistic concurrency* control protocols - they assume transactions will conflict and take steps to avoid it. i.e., they address the concurrency issues *before while the transaction is executing and* before the transaction commits.
- 2PL and TO are also *syntactic concurrency control protocols* as they deal only with the syntax (set of read and write operations) of the transactions.

Optimistic Concurrency Control (2)

- In an *optimistic concurrency control protocol*, we assume that most of the time, transactions will not conflict thus all of the locking and timestamp checking are not necessary.
- No checking for serialization is done while the transaction is executing
- During transaction execution, all updates are applied to *local copies* of the data items that are kept for the transaction
- During a validation phase the transactions updates are check to see if they violate serializability

Basic idea of OCC

- The idea behind OCC is to do all the checks at once
- If there is little interference between transactions, most will be validated successfully.
- Extra requirements for OCC:
 - Local Copy
 - Transaction Timestamps
 - Must keep track of write_set & read_set

Optimistic Concurrency Control in three stages

1. **Read Stage:** Transactions can read any data item. Writes are done to a local copy of the data item e.g., recorded in a log.
2. **Validation stage:** Transactions containing Write operations that are about to commit are *validated to see if the schedule meets the serializability* requirements.
3. **Write stage:** If the transaction will not conflict with other transactions, then it will be committed (writes to local copy applied to the database). Otherwise, the transaction will be rolled back.

RECAP

- Concurrency control motivation
 - If we insist only one transaction can execute at a time, in serial order, then performance will be quite poor.
- **Concurrency Control is a method for controlling or scheduling the operations of transactions in such a way that concurrent transactions can be executed safely** (i.e., without causing the database to reach an inconsistent state)
- If we do concurrency control properly, then we can maximize transaction throughput while avoiding any chance of corrupting the database