

**CMPT 506**

# **Database Concurrency Control**



**Read Chapters 21 and 22**

**Dr. Abdelkarim Erradi**

Department of Computer Science and Engineering

**QU**

# Outline

- Transaction Management
- Concurrency control
- Lock-based Concurrency control
- Transaction Isolation Levels
- Time-based Concurrency control
- Optimistic Concurrency Control

# Transaction Management

# Concept of Transaction

- Transaction = Logical unit of work on the database
  - Transfer money from one bank account to another
  - Checkout: place order and process payment
- A transaction consists of a sequence of read / write operations that must be performed as a **single** logical unit that must either commit or abort
  - Transaction boundaries are defined by the database user / application programmer
- **Atomicity, Consistency and Isolation are achieved using DB Transactions**

# Flight Reservation Example

## //1. Read customer info

```
SELECT customer_id, customer_lname, customer_fname  
INTO :v_customer_id, :v_customer_lname, :v_customer_fname  
FROM customer WHERE customer_id = :v_input_cid;
```

## //2. Write reservation information


```
UPDATE flight_reservations  
SET seat_assignment = :v_customer_id  
WHERE airline = 'ACME' AND flight_number = 123  
AND departure_date = '05-JAN-2017'  
AND departure_city = 'Dubai';
```

## //3. Write charges

```
INSERT INTO credit_card_charges  
VALUES (sysdatetime, :v_customer_id, :v_customer_lname,  
:v_customer_fname, 'VISA', :credit_card_number,  
:exp_date);
```

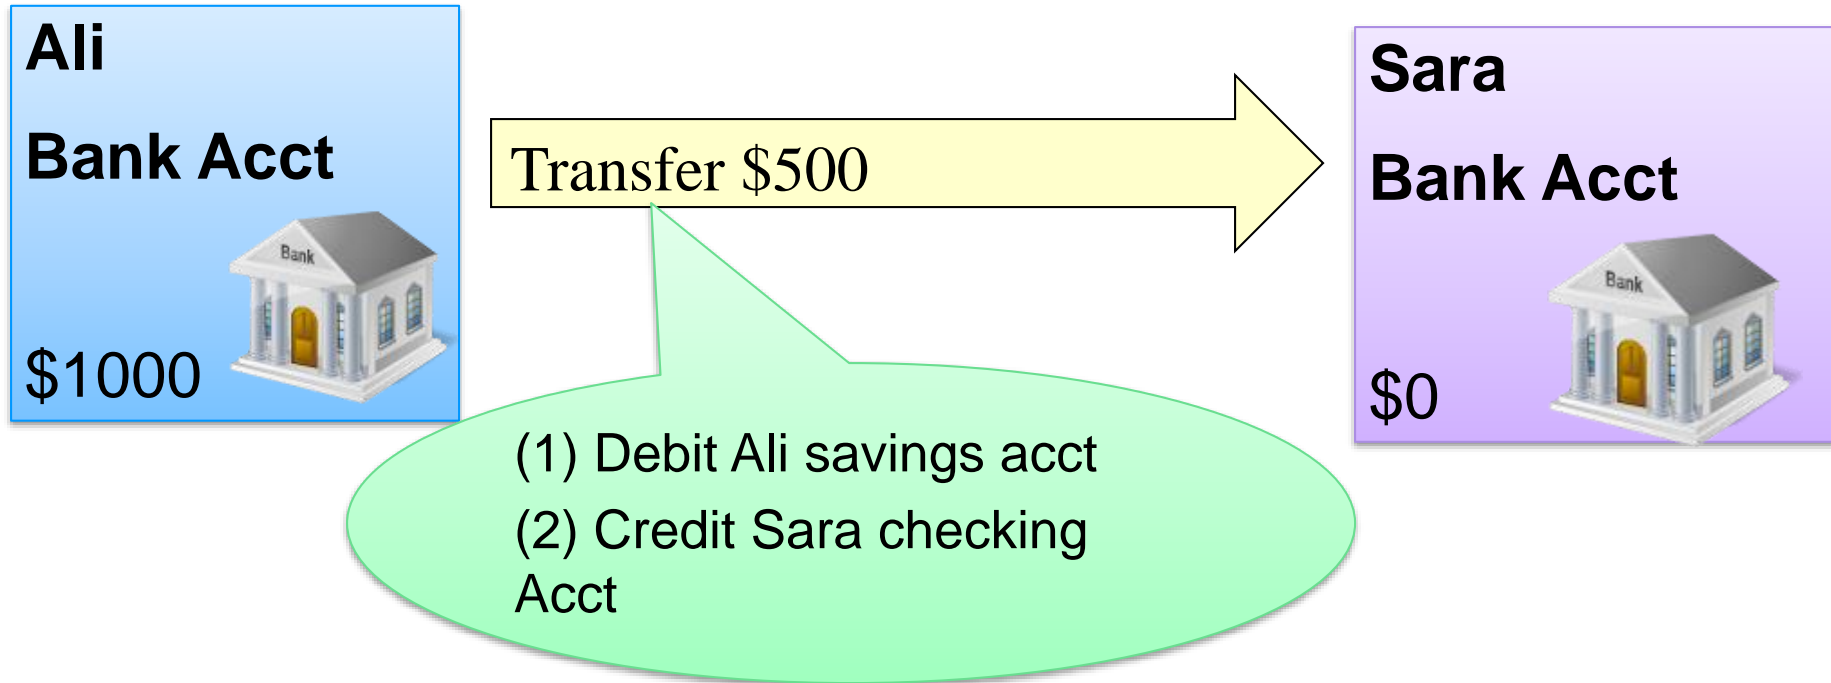
# Atomicity

- Consider a bank transaction T:



- T is transferring \$100 from B's account to A's account.
  - What if there is an error right after the first statement of T has been executed, i.e. the second statement is not executed?
- => You will get **partial update** = inconsistent state
- The DBMS has to ensure that every transaction is treated as an *atomic* unit, i.e. either all succeed otherwise rollback

# A transaction is a sequence of operations that must be executed as a whole



**Either both (1) and (2) happen or neither!**

**Every DB action takes place inside a transaction**

# We abstract away most of the application code when thinking about transactions

**User's point of view**

Transfer \$500

1. Debit savings
2. Credit checking

**Programmer's point of view**

Read Balance1  
Write Balance1  
Read Balance2  
Write Balance2

**DB's point of view**

Transaction = a sequence of DB reads (R) and writes (W)

T: R(A), W(A), R(B), W(B)

→ time



# SQL Transaction

- By default, each SQL statement (any query or modification of the database or its schema) is treated as a separate transaction
- Transactions can also be defined explicitly
  - `START TRANSACTION;`
  - `<sequence of SQL statements>`
  - `COMMIT;`
  - or `ROLLBACK;`
- COMMIT makes all modifications of the transaction permanent, ROLLBACK undoes all DB modifications made

# Transactions **ACID** properties

- **Atomic:** “all or nothing”, a transaction is performed entirely or not at all. Actions are never left partially executed.
- **Consistent:** Application level correctness => Data affected meet all validation rules such as constraints
- **Isolated:** Transactions should be independent from all other concurrent transactions and not interfere with each other (a.k.a. serializability) => e.g., The updates of a transaction must not be made visible to other transactions until it is **committed**
- **Durable:** Permanent side-effects upon commit => Results from completed transactions survive failures

# Concurrency control

# Concurrency Control

- Multiple **concurrent** transactions  $T_1, T_2, \dots$  may read/write Data Items  $A_1, A_2, \dots$  concurrently
- Concurrency Control is the process of managing concurrent operations performed on shared data so that data manipulation does not generate inconsistent databases or produce wrong results
- Objective
  - Maximise throughput (i.e., work performed)
  - Minimize response time
- Constraint
  - Avoid **interference** between transactions

# Three Concurrency Anomalies

- **Lost update** (some changes to DB get overwritten)
  - Two transactions  $T_1$  and  $T_2$  both modify the same data
  - $T_1$  and  $T_2$  both commit
  - Final state shows effects of only  $T_1$  but not of  $T_2$
- **Dirty read**
  - $T_1$  reads data written by  $T_2$  while  $T_2$  has not committed
  - If  $T_2$  aborts then  $T_1$  will have dirty data
- **Unrepeatable read**
  - Getting inconsistent results when a read operation is re-executed within a Transaction  $T$

# Illustrative Example

- **Example (to illustrate consistency issues that can be introduced by concurrent updates)**
- Ali at ATM1 withdraws \$100
- Sara at ATM2 withdraws \$50
- Initial balance = \$400, final balance = ?
  - Should be \$250 no matter who goes first

Read balance from DB;

If balance > withdrawalAmount {  
    balance = balance - withdrawalAmount;  
    Write balance to DB;

}

# No concurrent transactions scenario

Ali withdraws \$100:

```
read balance; ==> $400
if balance > amount then
    balance = balance - amount; ==> $300
write balance; ==> $300
```

Sara withdraws \$50:

```
read balance; ==> $300
if balance > amount then
    balance = balance - amount; ==> $250
write balance; ==> $250
```

# Lost update problem

Ali withdraws \$100:

```
read balance; => $400
```

Sara withdraws \$50:

```
read balance; => $400
```

```
If balance > amount then
```

```
balance = balance - amount; => $350
```

```
write balance; => $350
```

```
if balance > amount then
```

```
balance = balance - amount; => $300
```

```
write balance; => $300
```



**Lost update problem => DB is in inconsistent state**



# Lost update problem

Ali withdraws \$100:

```
read balance; => $400
```

```
if balance > amount then  
  balance = balance - amount; => $300  
write balance; => $300
```

Sara withdraws \$50:

```
read balance; => $400
```

```
if balance > amount then  
  balance = balance - amount; => $350  
write balance; => $350
```




**Lost update problem => DB is in inconsistent state**

# Dirty read problem

What will be the final account balance?



## Transaction 1:

Add \$100 to  
account A



## Transaction 2:

Add \$200 to  
account A

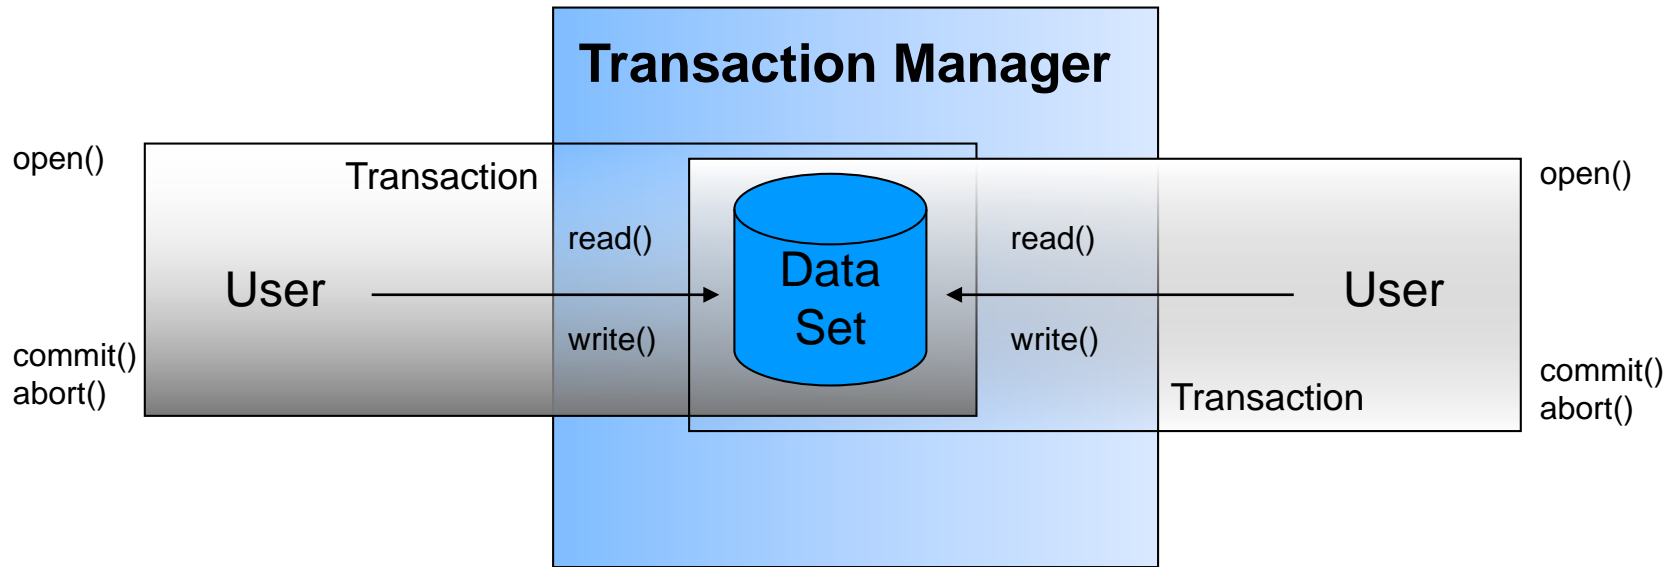


**Dirty read problem**

# Isolation

- To preserve consistency the DBMS must manage concurrency to **guarantee the Isolation property**.
  - Independence from all other transactions (serializability) => **same results as if the statements would have been executed in a single user scenario**
  - A transaction should not be affected by other concurrently running transactions

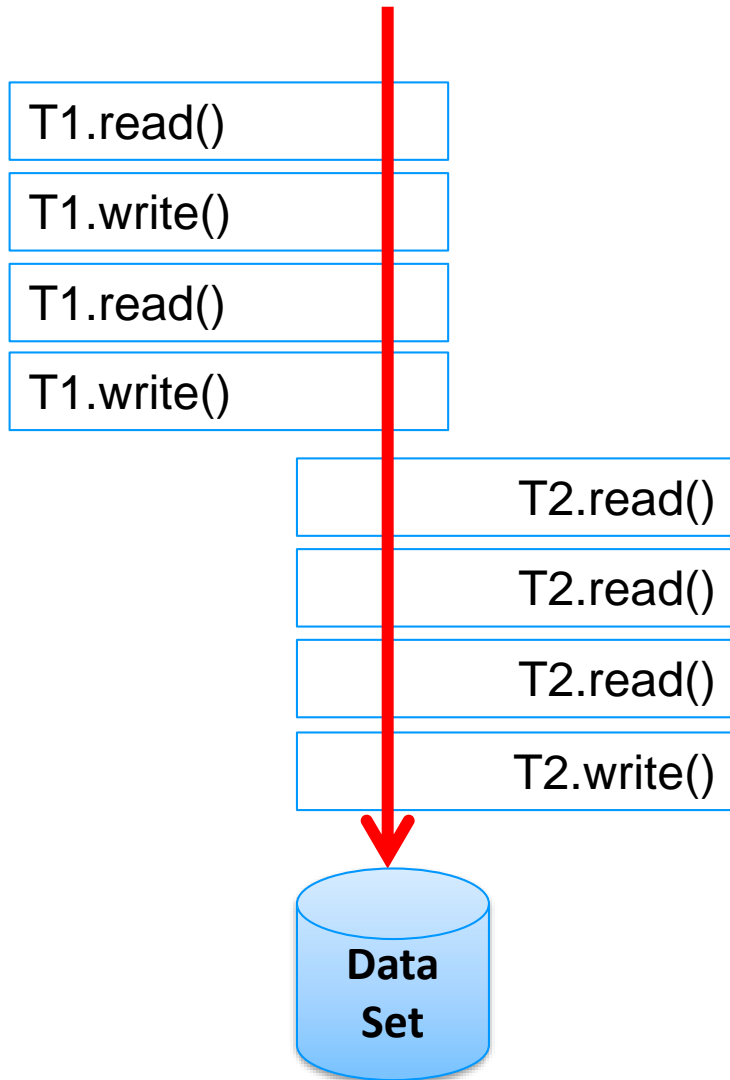
# Scheduling Concurrent Transactions



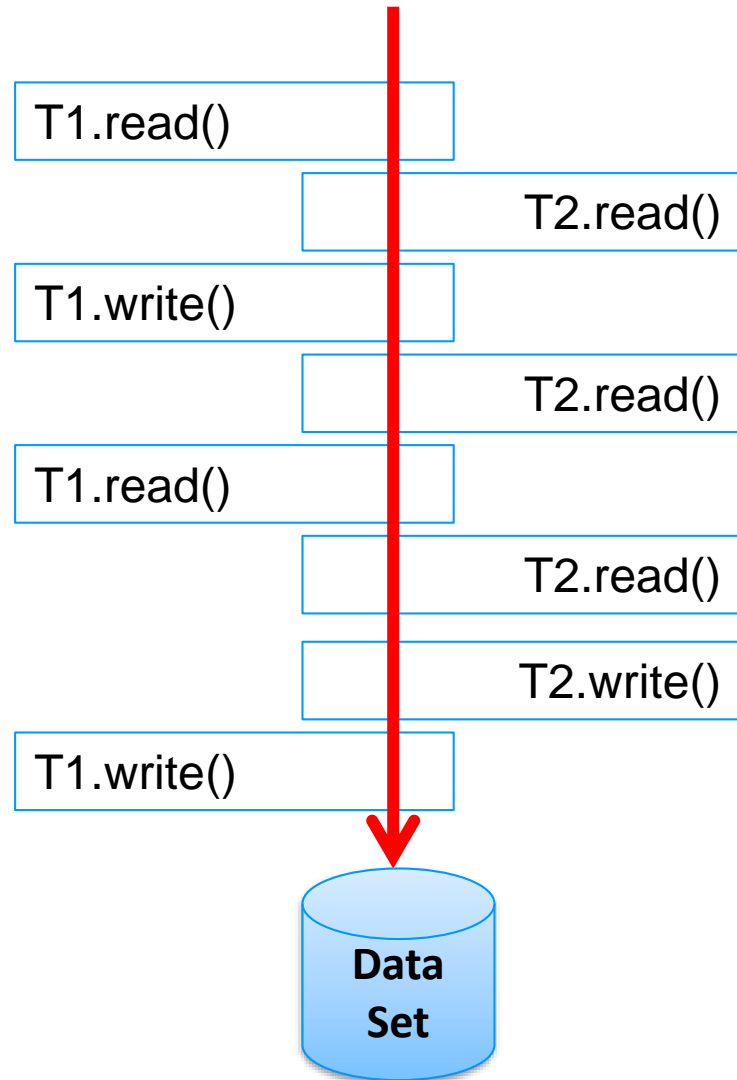
- The transaction scheduler has to organise or “schedule” database read and write actions of all concurrent transactions
- A “**schedule**” is a sequence of operations from different transactions that may be executed in an interleaved fashion
  - Such a schedule may compromise the integrity / consistency of a database

# Serial vs. Non-serial Schedule

## Serial Schedule



## Non-serial / Interleaved Schedule



# Serial schedule in which $T_1$ precedes $T_2$

$T_1$	$T_2$	$A$	$B$
		25	25
READ(A,t)			
t := t+100			
WRITE(A,t)		125	
READ(B,t)			
t := t+100			
WRITE(B,t)			125
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		250

Constraint of  $A = B$  is satisfied after this schedule

# Serial schedule in which $T_2$ precedes $T_1$

$T_1$	$T_2$	$A$	$B$
		25	25
	READ(A, s)		
	s := s*2		
	WRITE(A, s)	50	
	READ(B, s)		
	s := s*2		
	WRITE(B, s)		50
READ(A, t)			
t := t+100			
WRITE(A, t)		150	
READ(B, t)			
t := t+100			
WRITE(B, t)			150

Constraint of  $A = B$  is satisfied after this schedule

# Interleaved Transaction Schedules

- We want that the database system schedules transactions in an interleaved fashion:
  - Improve the responsiveness and increase throughput
- Interleaved schedules also create problems
  - Transactions may overwrite each others' updates
  - Transactions may base their calculations on retrieved data that is already out-of-date or on “dirty reads”



# Serial Transaction Schedules

- In order to avoid the concurrency problems described, one obvious solution would be to **schedule only one transaction at a time** for execution
- Such a completely “serialised” schedule will ensure that the **transactions are completely isolated** and cannot interfere with each other
  - But this strategy will reduce **concurrency** and **throughput**

# Conflict-Serializable Schedule

- DB will try to find a non-serial schedule that is ***equivalent*** to a serial schedule:
  - Schedule is **conflict-serializable** if its effect on the database state is the same as that of some serial schedule
- A **conflict-serializable schedule** is used by the Transaction Manager to schedule operations of different transactions in a way that interference and problems such as “lost updates” are avoided

# *Read and write operation conflict rules*

<i>Operations of different transactions</i>		<i>Conflict</i>	<i>Reason</i>
<i>read</i>	<i>read</i>	No	Because the effect of a pair of <i>read</i> operations does not depend on the order in which they are executed
<i>read</i>	<i>write</i>	Yes	Because the effect of a <i>read</i> and a <i>write</i> operation depends on the order of their execution
<i>write</i>	<i>write</i>	Yes	Because the effect of a pair of <i>write</i> operations depends on the order of their execution

# A serializable, but not serial, schedule

$T_1$	$T_2$	$A$	$B$
		25	25
READ(A,t)			
t := t+100			
WRITE(A,t)		125	
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
READ(B,t)			
t := t+100			
WRITE(B,t)			125
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		250

Constraint of  $A = B$  is satisfied after this schedule

# Non-serializable schedule

$T_1$	$T_2$	$A$	$B$
		25	25
READ(A,t)			
t := t+100			
WRITE(A,t)		125	
	READ(A,s)		
	s := s*2		
	WRITE(A,s)	250	
	READ(B,s)		
	s := s*2		
	WRITE(B,s)		50
READ(B,t)			
t := t+100			
WRITE(B,t)			150

Constraint of  $A = B$  is violated after this schedule => Non-serializable schedule


# Verifying Conflict-Serializability using a Precedence Graph?

- Precedence Graph = {**Nodes**: transactions, **Arcs**: r/w or w/w conflicts}
- The precedence graph for a schedule  $S$  contains:
  - A node for each committed transaction in  $S$
  - An **arc** from  $T_i$  to  $T_j$  if an action of  $T_i$  **precedes and conflicts with** one of  $T_j$ 's actions.
- The **schedule  $S$  is serializable if and only if the precedence graph has no cycles.**

# Example 1

---

T1	T2	T3
Read(A)		
...		
Write(A)		
		Read(A)
		...
		Write(A)
	Read(B)	
	...	
	Write(B)	
Read(B)		
...		
Write(B)		



**serial execution?**


Conflict-serializable Schedule: **T2, T1, T3**

(Notice that T3 should go after T2, although it starts before it!) => always 'read before write'

# Precedence Graph Example 2

$$D = \begin{bmatrix} T1 & T2 & T3 \\ R(A) & & \\ & W(A) & \\ W(A) & & \\ & & W(A) \end{bmatrix}$$

$$D = R_1(A) \ W_2(A) \ W_1(A) \ W_3(A)$$



**As  $T_1$  and  $T_2$  constitute a cycle the above schedule is not conflict-serializable.**



# Example 3 - 'Lost-update' problem

T1  
Read(N)

$N = N - 1$


Write(N)

T2

Read(N)

$N = N - 1$

Write(N)



Cycle  $\rightarrow$  not conflict-serializable

Not equivalent to any serial execution (why not?)

**We can draw a precedence graph to prove this**

*Here,  $T_1$  changes  $N$ , hence  $T_2$  should have either run first (read and write) or after (reading the changed value)*

# RECAP

- Concurrency control motivation
  - If we insist only one transaction can execute at a time, in serial order, then performance will be poor.
- **Concurrency Control is a method for controlling or scheduling the operations of transactions in such a way that concurrent transactions can be executed safely** (i.e., without causing the database to reach an inconsistent state)
- If we do concurrency control properly, then we can maximize transaction throughput while avoiding any chance of corrupting the database

# Concurrency Control Protocols

- The basis of concurrency control is protocols to maintain serialization in DBMSs
- E.g. protocols
  - **Two-Phase Locking (2PL)**
  - **Timestamps Ordering**
  - **Optimistic Concurrency Control (OCC)**
  - **Etc.**

# Lock-based Concurrency control

# Enforcing Serializability by Locks


**How to achieve correct concurrency control?**

- **Locks** = most popular solution for concurrency control
- Lock manager: grants/denies lock requests
- Locking mechanisms prevent conflicts
  - Readers block writers
  - Writers block readers

# How lock works

- Transaction needs lock before read or write
- Transaction must ask for lock
- If it does not get the lock
  - maybe another transaction already holds the lock
    - it is suspended
    - put on queue waiting for lock
- When transaction releases lock
  - Some suspended transaction awakened and given lock.

# Lock State Diagram




# Multi-transaction Rules


- Multiple Read Locks OK
  - Called Shared Lock
  - Several transactions can read same info
- Write Lock must be only lock on item
  - 2 Write Locks lead to race condition: lost update
  - Sharing with Read lock leads to dirty read or non-repeatable read.



# Read Lock State Diagram



# Write Lock State Diagram



# ReadLock WriteLock Compatibility Matrix

First  
Transaction  
Holds

Read Lock

Write Lock

Second Transaction Wants

Read Lock



Write Lock




# Locking protocol

- A **locking protocol** is a set of rules followed by all transactions while requesting and releasing locks
- Most popular protocol:
  - Two-phase locking (2PL), a transaction is not allowed to acquire any new locks after it has released a lock
  - **Strict 2 Phase Locking (S2PL)** = a transaction must hold all its exclusive locks till it commits/aborts
  - **Rigorous two-phase locking (R2PL)** is even stricter: here *all* locks are held till commit/abort
- **THEOREM**: if all transactions obey 2PL -> all schedules are serializable

# 2PL

- Phase 1: **Growing Phase**
  - transaction may obtain locks
  - transaction may not release locks
- Phase 2: **Shrinking Phase**
  - transactions issue no lock/upgrade request, after the first unlock/downgrade
  - transaction may release locks
  - transaction may not obtain locks
- 2PL assures serializability

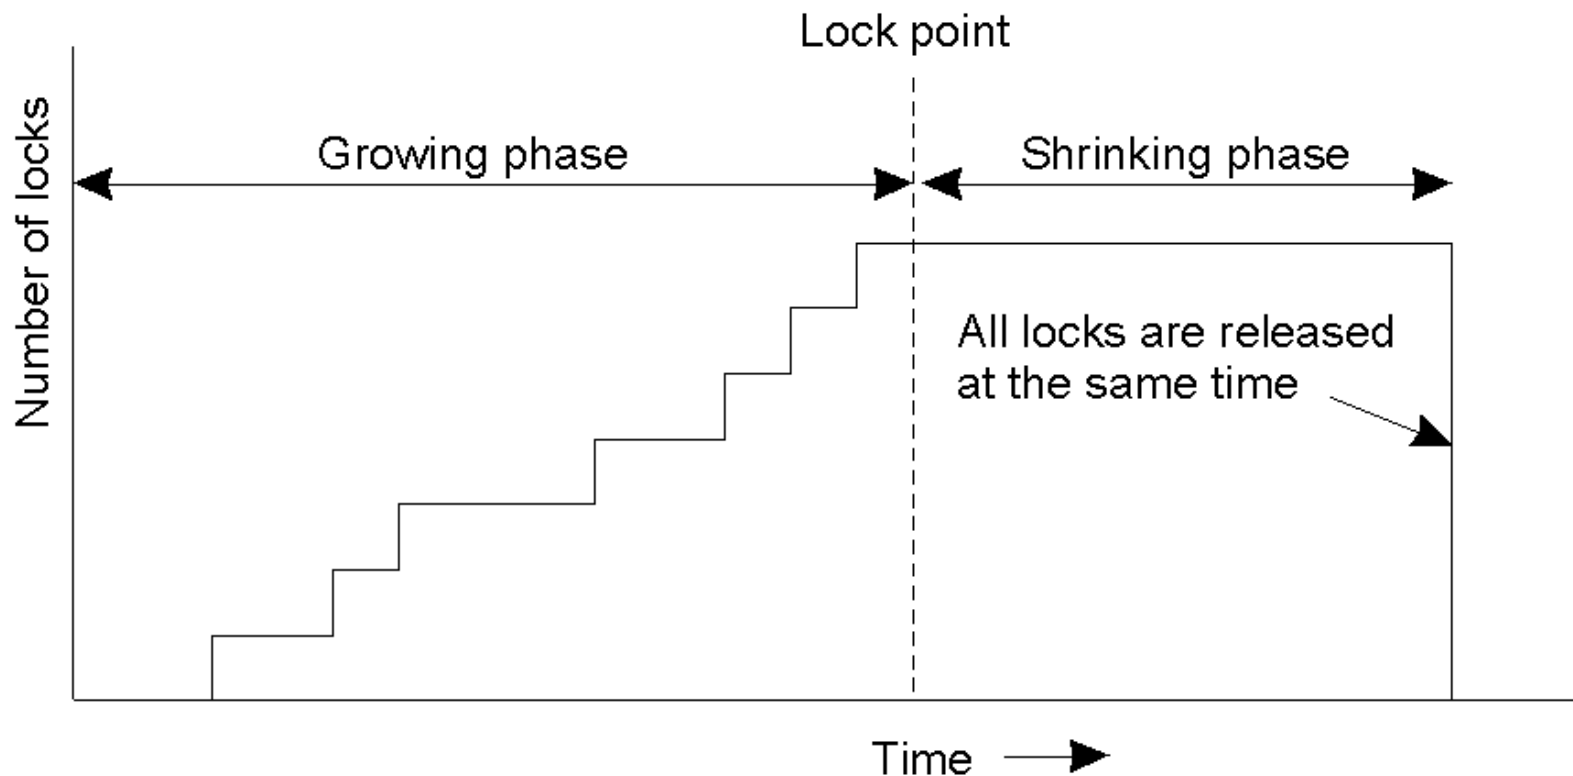
# Two-Phase Locking (1)



In two-phase locking, a transaction is not allowed to acquire any new locks after it has released a lock

# Strict Two-Phase Locking (2)

- Strict two-phase locking.



## Strict 2PL

T1	T2
s-lock(A)	
read(A)	
	s-lock(A)
x-lock(B)	
unlock(A)	
read(B)	
write(B)	
	read(A)
	unlock(A)
commit	
unlock(B)	
	s-lock(B)
	read(B)
	unlock(B)
	commit

Vs.

## Rigorous 2PL

T1	T2
s-lock(A)	
read(A)	
	s-lock(A)
x-lock(B)	
	read(A)
read(B)	
write(B)	
commit	
unlock(B)	
	s-lock(B)
	read(B)
unlock(A)	
	commit
	unlock(A)
	unlock(B)



# Use of locks in strict two-phase locking

1. When an operation accesses an object within a transaction:
  - (a) If the object is not already locked, it is locked and the operation proceeds.
  - (b) If the object has a conflicting lock set by another transaction, the transaction must wait until it is unlocked.
  - (c) If the object has a non-conflicting lock set by another transaction, the lock is shared and the operation proceeds.
  - (d) If the object has already been locked in the same transaction, the lock will be promoted if necessary and the operation proceeds. (Where promotion is prevented by a conflicting lock, rule (b) is used.)
2. When a transaction is committed or aborted, the server unlocks all objects it locked for the transaction.


# Why does strict 2PL imply serializability?

- Suppose that  $T'$  will perform an operation that conflicts with an operation that  $T$  has done:
  - $T'$  will update data item  $X$  that  $T$  read or updated
- $T$  must have had a lock on  $X$  that conflicts with the lock that  $T'$  wants
- $T$  won't release it until it commits or aborts
- So  $T'$  will wait until  $T$  commits or aborts

# Deadlocks

- The use of locks may lead to deadlocks
- **Definition**
  - A deadlock is a state in which each member of a group of transactions is waiting for some other member to release a lock
- The scheduler is responsible for detecting and breaking the deadlock
  - Deadlocks may be broken by simply aborting one of the transactions involved
- How do you choose which transaction to abort?
  - Abort the oldest
  - Abort depending on the complexity of the transactions

# Deadlocks



- Auto-detection and “resolution”

Msg 1205, Level 13, State 51, Line 1

Transaction (Process ID 232) was deadlocked on lock resources with another process and has been chosen as the deadlock victim. Rerun the transaction

# Deadlock Example

## Transaction A

Read X **Writelock X**

Read Y **Readlock Y**

Write X

Cannot get lock  
due to B's  
Writelock -- A  
sleeps

**Each sleeps, waiting for the other**

## Transaction B

**Writelock Y**  
Read Y

**Readlock X**  
Read X

Write Y

Cannot get  
lock due to  
A's Writelock  
-- B sleeps

# Handling Deadlocks

- Deadlock Detection
  - Use the Precedence graph to identify cycles and select transactions to be aborted
    - Select the transaction involved in most of the cycles
    - Select the oldest transactions
    - Select the one that did the least amount of work
- Deadlock Prevention
  - Lock all objects at the very beginning of a transaction in one atomic action
  - Problem
    - Reduced concurrency: unnecessary access restriction to shared resources

# Handling Deadlocks using Timeout

- Rather than detecting deadlock (an overhead), you could just use timeouts, but how long should the timeout be?
- Timeouts
  - Each lock is given a period of time where it is invulnerable
  - After this timeout, it becomes vulnerable
  - If transaction X holds a lock that becomes vulnerable and transaction Y is waiting for X, then X is aborted
  - Problem
    - Hard to decide on an appropriate length of timeout
    - Transactions may be aborted, when the lock becomes vulnerable and another transaction waits, but there is no deadlock

# Transaction Isolation Levels



# Transaction Isolation Levels


- Determine how much one transaction impacts another
- 4 standard levels
  - Read Uncommitted
  - Read Committed (default in many DMBS)
  - Repeatable Read
  - Serializable
- The *isolation level* of a transaction defines what data *that* transaction may see

# Read Uncommitted

- Does not take or honour locks
- Can read uncommitted data
  - *Dirty data* is data that has been modified by a transaction that has not yet committed.
  - If that transaction is rolled back after another transaction has read its dirty data, inconsistency is introduced
- Sacrificing consistency in favour of high concurrency
- Useful for reporting applications
- Be very careful!

# Read Uncommitted

- Read Uncommitted
  - Dirty reads might occur



## Transaction 1

```
UPDATE EMPLOYEE SET SALARY = 1300  
WHERE EMP_ID = 123
```

```
ABORT;
```

## Transaction 2


```
SELECT * FROM  
EMPLOYEE  
WHERE SALARY > 1000
```

# Read Committed (the default)

- Cannot read uncommitted data
- Share lock before reading data and released after processing is complete
- Dirty reads do not happen but non-repeatable reads can happen

# Read Committed

- Read Committed
  - Non-repeatable reads might occur



## Transaction 1

```
SELECT * FROM  
  EMPLOYEE  
WHERE EMP_ID = 123
```

```
SELECT * FROM  
  EMPLOYEE  
WHERE EMP_ID = 123
```

## Transaction 2


```
UPDATE EMPLOYEE SET SALARY = 1300  
  WHERE EMP_ID = 123;  
COMMIT;
```

# Repeatable Read

- All **data** records read by a SELECT statement cannot be changed
- Holds shared locks on data read until transaction commit/rollback
- Rows that have been read can be read again with confidence they won't have changed
  - A query running more than once within the same transaction returns same values
- If the SELECT statement contains any ranged WHERE clauses, phantom reads can occur
- Reduce concurrency and degrade performance

# Repeatable Reads

- Repeatable Reads
  - Phantom Reads might occur



## Transaction 1

```
SELECT * FROM  
EMPLOYEE  
WHERE SALARY > 1000
```

```
SELECT * FROM  
EMPLOYEE  
WHERE SALARY > 1000
```

## Transaction 2

```
INSERT INTO EMPLOYEE (EMP_ID, NAME, SALARY)  
VALUES("123", "Anil", 1200);  
COMMIT;
```

# Phantom Problem

- A “phantom” is a tuple that is invisible during part of a transaction execution but appears when the query is re-executed
- In our example:
  - T1: reads list of products
  - T2: inserts a new product
  - T1: re-reads: a new product appears !



# Dealing With Phantoms

- Lock the entire table, or
- Lock the index entries of qualifying rows
  - if index is available

Dealing with phantoms is expensive !

# Serializable

- This isolation level specifies that all transactions occur in a completely isolated fashion
  - *i.e.*, as if all transactions in the system had executed serially, one after the other
- Locks index ranges as well as rows or table locks
- Phantom rows will not appear if the same query is issued twice within a transaction
- Greatly reduces concurrency

# Concurrency and Consistency

Isolation Table	Dirty Read	Non-Repeatable Read	Phantom Read
Read Uncommitted	Possible	Possible	Possible
Read Committed	Not Possible	Possible	Possible
Repeatable Read	Not Possible	Not Possible	Possible
Serializable	Not Possible	Not Possible	Not Possible

- Concurrency and consistency are mutually opposing goals

# Time-based Concurrency control

# Timestamp Ordering

- The timestamp of a transaction  $T$  is the time at which that transaction was initiated in the DBMS:  $TS(T)$
- We can use clock time or an incremental identifier (counter) for  $TS(T)$
- Two timestamps are also associated with each data item  $x$ .
  1.  $read\_TS(x)$  is the  $TS(T)$  of the last transaction  $T$  to read from  $x$ .
  2.  $write\_TS(x)$  is the  $TS(T)$  of the last transaction  $T$  to write to  $x$ .

# Protocol Rules

- Two simple rules to follow:
  1. Before T issues a write(x), check to see if
    - $TS(T) < read\_TS(x)$  or if
    - $TS(T) < write\_TS(x)$
    - If so, then abort transaction T
    - If not, then perform write(x) and set  $write\_TS(x) = TS(T)$
  2. Before T issues a read(x), check to see if
    - $TS(T) < write\_TS(x)$  Then abort transaction T.
    - if  $TS(T) \geq write\_TS(x)$  then execute read(x) and set  $read\_TS(x) = TS(T)$  only if  $TS(T)$  is greater than the current  $read\_TS(x)$

# Advantages and Limitation

- When a transaction is aborted, it is the restarted and issued a new  $TS(T)$ .
- Note that with timestamp ordering, deadlock can not occur.
- However, starvation is possible i.e., a transaction keeps getting aborted over and over.

# Cascading rollbacks

- Timestamp ordering can also produce cascading rollbacks:
  - Assume transaction T begins executing and performs some read and write operations on data items *a*, *b* and *c*
  - However, T then reaches a data item it can not read or write and T must then be aborted.
  - Any effects of transaction T must then be rolled back.
  - Before T aborts, however, other transactions (T1, T2 and T3) have read and written data items *a*, *b* and *c* so *these other transactions must also be rolled back*.
  - There may be other transactions (T4 and T5) that worked with data items read or written by T1, T2 and T3, etc.



# Optimistic Concurrency Control

# Pessimistic Concurrency Control

- Two Phase Locking (2PL) and Timestamp Ordering (TO) are *pessimistic concurrency control* protocols:
  - They assume transactions will conflict and take steps to avoid it. i.e., they address the concurrency issues *while the transaction is executing and* before the transaction commits

# Optimistic Concurrency Control

- In an *optimistic concurrency control protocol*, we assume that most of the time, transactions will not conflict thus all of the locking and timestamp checking are not necessary while the transaction is executing
- During transaction execution, all updates are applied to *local copies* of the data items that are kept for the transaction
- During a validation phase the transactions updates are check to see if they violate serializability

# Optimistic Concurrency Control in three stages

1. **Working Stage:** Transactions can read any data item. Writes are done to a local copy of the data item e.g., recorded in a log.
2. **Validation stage:** Transactions containing Write operations that are about to commit are *validated to see if the schedule meets the serializability* requirements.
3. **Write stage:** If the transaction will not conflict with other transactions, then it will be committed (writes to local copy applied to the database). Otherwise, the transaction will be rolled back.

# Summary

- A transaction consists of a sequence of read / write operations that must be performed as a single logical unit
- The DBMS guarantees the atomicity, consistency, isolation and durability of transactions
- **Transactions manager** ensures that the database remains in a **consistent** (correct) state despite system failures (e.g., power failures and operating system crashes) and transaction failures
- **Concurrency-control manager** controls the interaction among the concurrent transactions, to ensure the consistency of the database (when multiple users concurrently update the same data)